

Deep Learning — Assignment 1

First assignment for the 2020 Deep Learning course (NWI-IMC058) of the Radboud University.

Gijs van Tulder (g.vantulder@cs.ru.nl) and Twan van Laarhoven (tvanlaarhoven@cs.ru.nl)

September 2020

Names: Ward Theunisse & Nienke Wessel

Group: 25

Instructions:

- Fill in your names and the name of your group.
- Answer the questions and complete the code where necessary.
- Re-run the whole notebook before you submit your work.
- Save the notebook as a PDF and submit that in Brightspace together with the `.ipynb` notebook file.

Objectives

In this assignment you will

1. Experiment with gradient descent optimization;
2. Derive and implement gradients for binary cross-entropy loss, the sigmoid function and a linear layer;
3. Test your gradient implementations with the finite differences method;
4. Use these components to implement and train a simple neural network.

```
In [1]: %matplotlib inline
import numpy as np
import scipy.optimize
import sklearn.datasets
import matplotlib.pyplot as plt

np.set_printoptions(suppress=True, precision=6, linewidth=200)
plt.style.use('ggplot')
```

1.1 Gradient descent optimization

Consider the following function with two parameters and its derivatives:


$$f(x, y) = x^2 + y^2 + x(y + 2) + \cos(3x)$$
$$\frac{\partial f}{\partial x} = 2x - 3 \sin(3x) + y + 2$$
$$\frac{\partial f}{\partial y} = x + 2y$$

```
In [2]: def f(x, y):
        return x ** 2 + y ** 2 + x * (y + 2) + np.cos(3 * x)
def grad_x_f(x, y):
    return 2 * x - 3 * np.sin(3 * x) + y + 2
def grad_y_f(x, y):
    return x + 2 * y
```

A plot of the function shows that it has multiple local minima:

```
In [3]: def plot_f_contours():
    xx, yy = np.meshgrid(np.linspace(-5, 5), np.linspace(-5, 5))
    zz = f(xx, yy)
    plt.contourf(xx, yy, zz, 50)
    plt.contour(xx, yy, zz, 50, alpha=0.2, colors='black', linestyle='solid')
    plt.xlabel('x')
    plt.ylabel('y')

plt.figure(figsize=(10, 7))
plot_f_contours()
```



Implement gradient descent

We would like to find the minimum of this function using gradient descent.

Implement the gradient descent updates for x and y in the function below:

```
In [4]: def optimize_f(x, y, step_size, steps):
    # keep track of the parameters we tried so far
    x_hist, y_hist = [x], [y]

    # run gradient descent for the number of steps
    for step in range(steps):
        # compute the gradients at the current point
        dx = grad_x_f(x, y)
        dy = grad_y_f(x, y)

        # apply the gradient descent updates to x and y
        x = x - step_size * dx # TODO: compute the update
        y = y - step_size * dy # TODO: compute the update

        # store the new parameters
        x_hist.append(x)
        y_hist.append(y)

    return x, y, f(x, y), x_hist, y_hist
```

Tune the parameters

We will now try if our optimization method works.

Use this helper function to plot the results:


```
In [5]: # helper function that plots the results of the gradient descent optimization
def plot_gradient_descent_results(x, y, val, x_hist, y_hist):
    # plot the path on the contour plot
    plt.figure(figsize=(20, 7))
    plt.subplot(1, 2, 1)
    plot_f_contours()
    plt.plot(x_hist, y_hist, '-.')

    # plot the learning curve
    plt.subplot(1, 2, 2)
    plt.plot(f(np.array(x_hist), np.array(y_hist)), '.r-')
    plt.title('Minimum value: %f' % f(x_hist[-1], y_hist[-1]))
```

Run the gradient descent optimization with the following initial settings:

$x=3$, $y=2$, $\text{step_size}=0.1$, $\text{steps}=10$

```
In [6]: results = optimize_f(x=3, y=2, step_size=0.1, steps=10)
        plot_gradient_descent_results(*results)
```

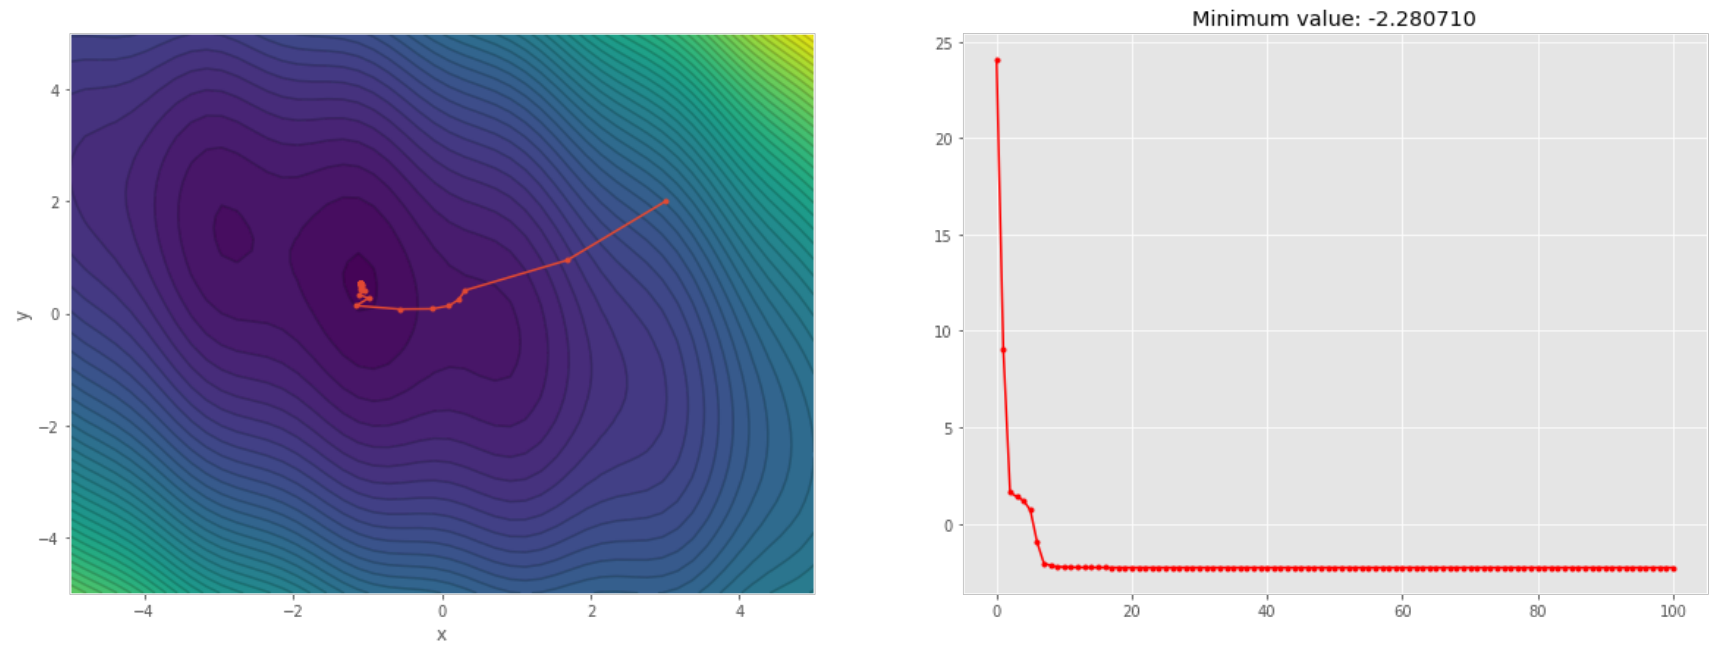


Does it find the minimum of the function? What happens?

The algorithm finds a local minimum but not the global minimum.

Try a few different values for the `step_size` and the number of `steps` to get closes to the optimal solution:

```
In [7]: # TODO: tune the parameters to find a better optimum
        results = optimize_f(x=3, y=2, step_size=0.15, steps=100)
        plot_gradient_descent_results(*results)
```



Were you able to find a step size that reached the global optimum? If not, why not?

Yes, we were. A smaller step size than 0.1 does not seem to help, but a slightly larger one (we used 0.15) 'skips' over the local minimum to the global minimum and then finishes there. (Too large a step size also does not work as it switches between different minima)

Implement a decreasing step size

You might get better results if you use a step size that is large at the beginning, but slowly decreases during the optimization.

Try the following scheme to compute the step size η_t in step t , given a decay parameter d :

$$\eta_t = \eta_0 d^t$$

Update your optimization function to use this step size schedule:

```
In [8]: def optimize_f(x, y, step_size, steps, decay=1.0):
# keep track of the parameters we tried so far
x_hist, y_hist = [x], [y]

# run gradient descent for the number of steps
for step in range(steps):
# compute the gradients at this point
dx = grad_x_f(x, y)
dy = grad_y_f(x, y)

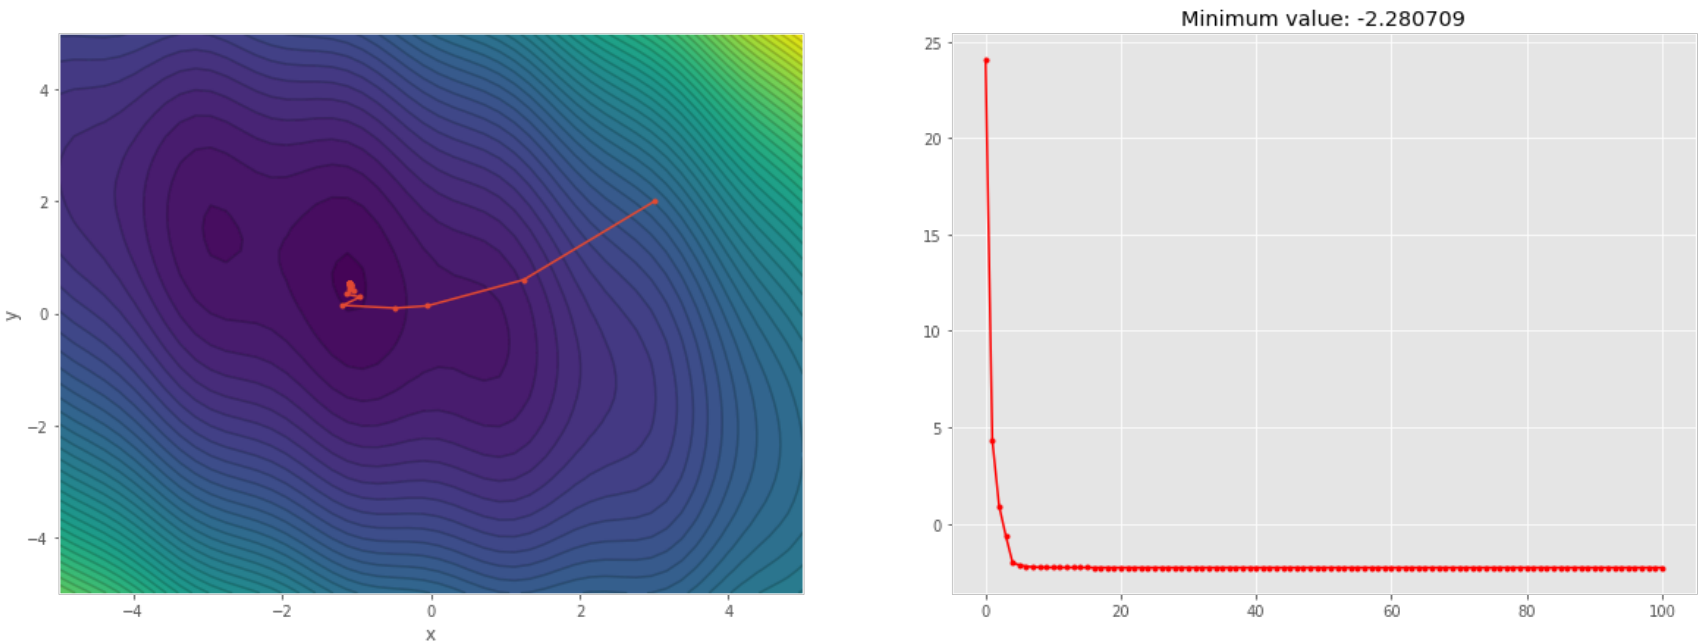
# apply the gradient descent updates to x and y
x = x - step_size * (decay ** step) * dx # TODO: compute the update including step size decay
y = y - step_size * (decay ** step) * dy # TODO: compute the update including step size decay

# store the new parameters
x_hist.append(x)
y_hist.append(y)

return x, y, f(x, y), x_hist, y_hist
```

Tune the `step_sizes` , `steps` and `decay` parameters to get closer to the global minimum:

```
In [9]: # TODO: tune the parameters to find the local optimum
results = optimize_f(x=3, y=2, step_size=0.2, steps=100, decay=0.95)
plot_gradient_descent_results(*results)
```



We will now look at some more complex functions that we can try to optimize.

1.2 Neural network components

In this assignment, we will implement a simple neural network from scratch. We need four components:


- 1. A sigmoid activation function,
- 2. A ReLU activation function,
- 3. A binary cross-entropy loss function,
- 4. A linear layer.

For each component, we will implement the forward pass, the backward pass, and the gradient descent update.

Sigmoid non-linearity

The sigmoid function is defined as:

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$



Give the derivative of the sigmoid function:

$$\begin{aligned} \frac{\partial \sigma(x)}{\partial x} &= \frac{\partial}{\partial x} (1 + e^{-x})^{-1} \\ &= -1 \cdot (1 + e^{-x})^{-2} \cdot (-e^{-x}) \\ &= (1 + e^{-x})^{-2} \cdot e^{-x} \\ &= (1 + e^{-x})^{-1} \cdot (1 + e^{-x})^{-1} \cdot e^{-x} \\ &= \sigma(x) \cdot (1 + e^{-x})^{-1} \cdot e^{-x} \\ &= \sigma(x) \cdot \frac{e^{-x}}{1 + e^{-x}} \\ &= \sigma(x) \cdot \frac{-1 + (1 + e^{-x})}{1 + e^{-x}} \\ &= \sigma(x) \cdot \left(\frac{-1}{1 + e^{-x}} + \frac{1 + e^{-x}}{1 + e^{-x}} \right) \\ &= \sigma(x) \cdot (-\sigma(x) + 1) \\ &= \sigma(x) \cdot (1 - \sigma(x)) \end{aligned}$$

Implement the sigmoid and its gradient in the functions `sigmoid(x)` and `sigmoid_grad(x)` :

```
In [10]: def sigmoid(x):
         return 1 / (1 + np.exp(-x))

         def sigmoid_grad(x):
             return sigmoid(x) * (1 - sigmoid(x))

         # try with a random input
         x = np.random.uniform(-10, 10, size=5)
         print('x:', x)
         print('sigmoid(x):', sigmoid(x))
         print('sigmoid_grad(x):', sigmoid_grad(x))

x: [ 5.044848 -1.422341 -8.288057 -1.842335  8.040757]
sigmoid(x): [0.993599 0.194295 0.000251 0.136775 0.999678]
sigmoid_grad(x): [0.00636  0.156544 0.000251 0.118068 0.000322]
```

To check that the gradient implementation is correct, we can compute the numerical derivative using the finite difference method. From [Chapter 11.5 of the Deep Learning book](http://www.deeplearningbook.org/contents/guidelines.html) (<http://www.deeplearningbook.org/contents/guidelines.html>):

Because

$$f'(x) = \lim_{\epsilon \rightarrow 0} \frac{f(x + \epsilon) - f(x)}{\epsilon},$$

we can approximate the derivative by using a small, finite ϵ :

$$f'(x) \approx \frac{f(x + \epsilon) - f(x)}{\epsilon}.$$

We can improve the accuracy of the approximation by using the centered difference:

$$f'(x) \approx \frac{f(x + \frac{1}{2}\epsilon) - f(x - \frac{1}{2}\epsilon)}{\epsilon}.$$

The perturbation size ϵ must be large enough to ensure that the perturbation is not rounded down too much by finite-precision numerical computations.

Use the central difference method to check your implementation of the sigmoid gradient. Compute the numerical gradient and check that it is close to the symbolic gradient computed by your implementation:

```
In [11]: # start with some random inputs
x = np.random.uniform(-2, 2, size=5)

# compute the symbolic gradient
print('Symbolic', sigmoid_grad(x))

# TODO: compute the numerical gradient
def num_gradient_sigmoid(x, epsilon):
    return (sigmoid(x + 0.5*epsilon) - sigmoid(x - 0.5*epsilon))/epsilon

print('Finite differences', num_gradient_sigmoid(x, 0.01))

Symbolic [0.139511 0.140679 0.105274 0.19541  0.246271]
Finite differences [0.139511 0.140679 0.105274 0.19541  0.246271]
```

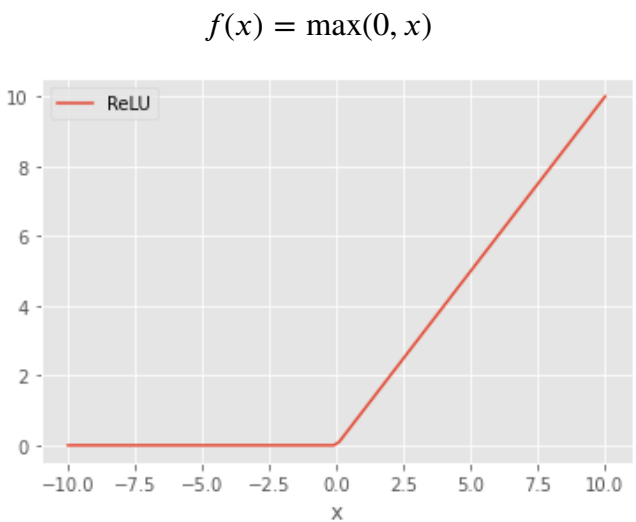
Is the gradient computed with finite differences exactly the same as the analytic answer? Why (not)?

They are not exactly the same, as the finite differences is only an approximation. In order for it to be exactly the same, one would have to make epsilon infinitely small, but that is not possible on a computer. With the current epsilon (0.01) they look the same for the six digits that are shown after rounding.

If there is a visible difference between the two gradients, please try to make this a small as possible before you continue.

Rectified linear units (ReLU)

The rectified linear unit is defined as:



Give the derivative of the ReLU function:

$$\begin{aligned} \frac{\partial f(x)}{\partial x} &= \frac{\partial}{\partial x} \max(0, x) \\ &= \frac{\partial}{\partial x} \begin{cases} 0, & \text{for } x < 0 \\ x, & \text{for } 0 \leq x \end{cases} \\ &= \begin{cases} \frac{\partial}{\partial x} 0, & \text{for } x < 0 \\ \frac{\partial}{\partial x} x, & \text{for } 0 \leq x \end{cases} \\ &= \begin{cases} 0, & \text{for } x < 0 \\ 1, & \text{for } 0 \leq x \end{cases} \end{aligned}$$

Strictly seen, ReLU is not differentiable at 0, but that is undesirable in a setting where we do need a result, so we've chosen to define it as part of the case where the max returns x.

Implement the ReLU function and its gradient in the functions `relu(x)` and `relu_grad(x)` . Use the finite differences method to check that the gradient is correct:


```
In [12]: def relu(x):
# DONE: implement the relu function
return np.maximum(0,x)

def relu_grad(x):
# DONE: implement the gradient of the relu function
return 0 if x < 0 else 1
relu_grad = np.vectorize(relu_grad)

# try with a random input
x = np.random.uniform(-10, 10, size=5)
print('x:', x)
print('relu(x):', relu(x))
print('relu_grad(x):', relu_grad(x))
print()

# DONE: compute and compare the symbolic and numerical gradients
print('Symbolic:', relu_grad(x))

def num_gradient_relu(x, epsilon):
return (relu(x + 0.5*epsilon) - relu(x - 0.5*epsilon))/epsilon

print('Finite differences', num_gradient_relu(x, 0.01))

x: [5.636688 7.488439 4.418439 6.663227 3.226479]
relu(x): [5.636688 7.488439 4.418439 6.663227 3.226479]
relu_grad(x): [1 1 1 1 1]

Symbolic: [1 1 1 1 1]
Finite differences [1. 1. 1. 1. 1.]
```

Comparing sigmoid and ReLU

The sigmoid and ReLU activation functions have slightly different characteristics.

Run the code below to plot the sigmoid and ReLU activation functions and their gradients:

```
In [13]: x = np.linspace(-10, 10, 100)

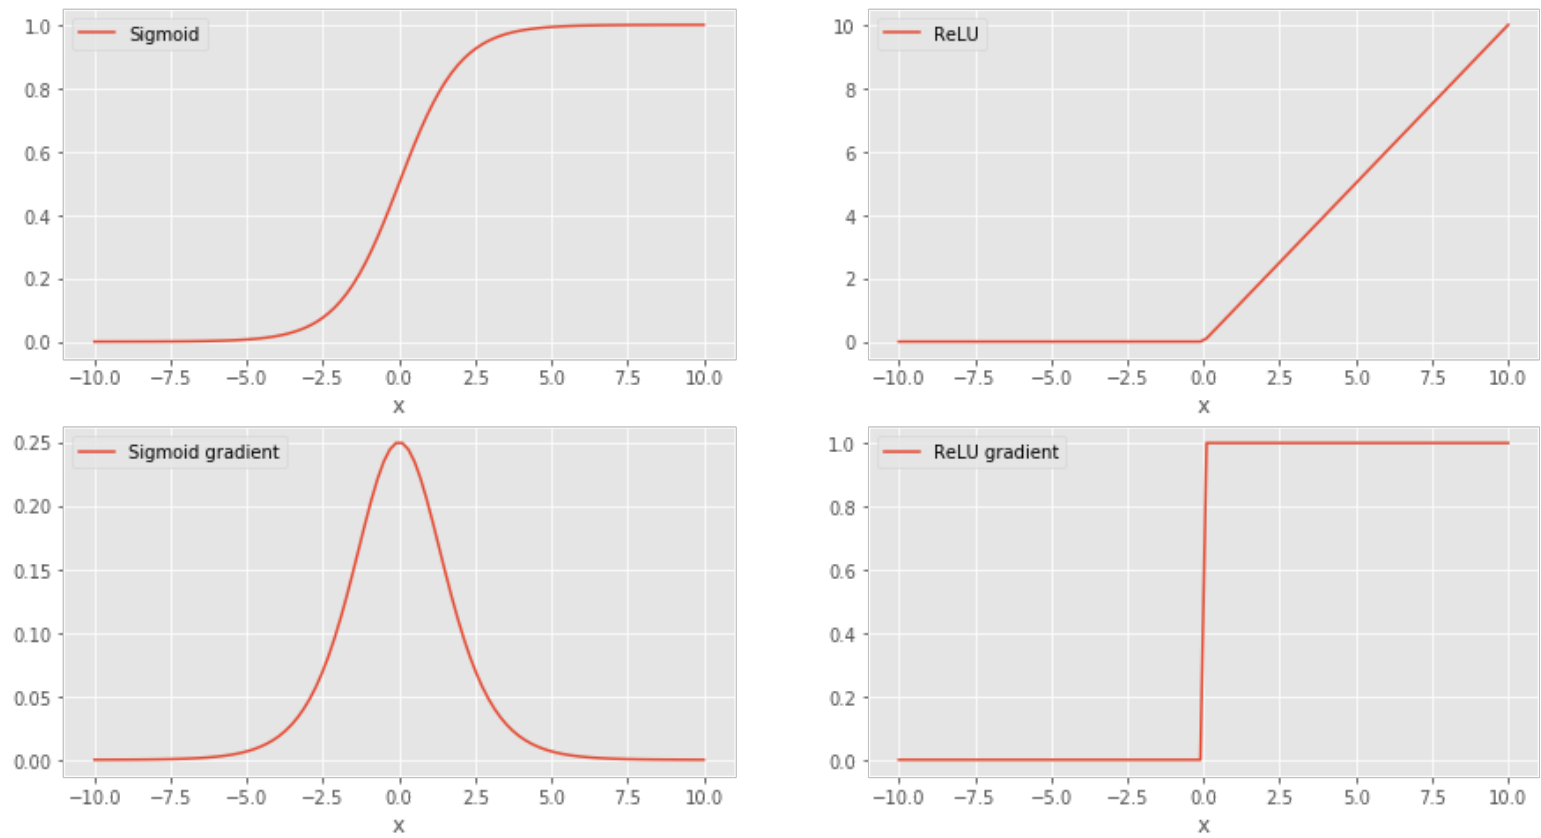
plt.figure(figsize=(15, 8))

plt.subplot(2, 2, 1)
plt.plot(x, sigmoid(x), label='Sigmoid')
plt.xlabel('x')
plt.legend(loc='upper left')

plt.subplot(2, 2, 2)
plt.plot(x, relu(x), label='ReLU')
plt.xlabel('x')
plt.legend(loc='upper left')

plt.subplot(2, 2, 3)
plt.plot(x, sigmoid_grad(x), label='Sigmoid gradient')
plt.xlabel('x')
plt.legend(loc='upper left')

plt.subplot(2, 2, 4)
plt.plot(x, relu_grad(x), label='ReLU gradient')
plt.xlabel('x')
plt.legend(loc='upper left');
```



Which activation function would you recommend for a network that outputs probabilities, i.e., outputs $\in (0, 1)$? Why?

Sigmoid. In the domain of probabilities, we're not just interested in the outer edges of the domain: we are also really interested in the probabilities that fall inbetween. W.r.t. the activation functions of the hidden layers, the sigmoid function also captures more granularity (more values of the state space are utilised).

Compare the gradients for sigmoid and ReLU. What are the advantages and disadvantages of each activation function?

ReLU is really easy / fast to compute, and the gradient is consistently large, even for very small values of x, so you don't suffer from a vanishing gradient (for positive values). However, it's gradient being zero for values below zero would tend it to not learn anymore and become "stuck", which is undesirable.

Sigmoid has a soft transition, leading to a larger granularit in outputs. However, for some values, the gradient approaches zero, which really slows learning.

Binary cross-entropy loss

We will use the binary cross-entropy loss to train our network. This loss function is useful for binary classification.

The binary cross-entropy (BCE) is a function of the ground truth label $y \in \{0, 1\}$ and the predicted label $\hat{y} \in (0, 1)$:

$$\mathcal{L} = -(y \log \hat{y} + (1 - y) \log(1 - \hat{y}))$$

To minimize the BCE loss with gradient descent, we need to compute the gradient with respect to the prediction \hat{y} .

Derive the gradient for the BCE loss:

$$\begin{aligned} \frac{\partial \mathcal{L}}{\partial \hat{y}} &= \frac{\partial}{\partial \hat{y}} -(y \log \hat{y} + (1 - y) \log(1 - \hat{y})) \\ &= -(\frac{\partial}{\partial \hat{y}}(y \log \hat{y}) + \frac{\partial}{\partial \hat{y}}((1 - y) \log(1 - \hat{y}))) \\ &= -(y * \frac{\partial}{\partial \hat{y}} \log \hat{y} + (1 - y) * \frac{\partial}{\partial \hat{y}} \log(1 - \hat{y})) \\ &= -(y * \frac{1}{\hat{y}} + (1 - y) * \frac{\partial}{\partial \hat{y}} \log(1 - \hat{y})) \\ &= -(\frac{y}{\hat{y}} - \frac{1 - y}{1 - \hat{y}}) \\ &= -(\frac{y(1 - \hat{y})}{\hat{y}(1 - \hat{y})} - \frac{\hat{y}(1 - y)}{\hat{y}(1 - \hat{y})}) \\ &= -(\frac{y - y\hat{y}}{\hat{y}(1 - \hat{y})} + \frac{-\hat{y} + \hat{y}y}{\hat{y}(1 - \hat{y})}) = -(\frac{y - \hat{y}}{\hat{y}(1 - \hat{y})}) = \frac{\hat{y} - y}{\hat{y}(1 - \hat{y})} \end{aligned}$$

Implement bce_loss(y, y_hat) and bce_loss_grad(y, y_hat) and use the finite differences method to check that the gradient is correct:

```
In [14]: def bce_loss(y, y_hat):
# DONE: implement the BCE loss
return -(y*np.log(y_hat) + (1-y) * np.log(1-y_hat))

def bce_loss_grad(y, y_hat):
# DONE: implement the gradient of the BCE loss
return (y_hat - y) / (y_hat * (1 - y_hat))

# try with some random inputs
y = np.random.randint(2, size=5)
y_hat = np.random.uniform(0, 1, size=5)
print('y:', y)
print('y_hat:', y_hat)
print('bceloss(y, y_hat):', bce_loss(y, y_hat))
print()

# TODO: compute and compare the symbolic and numerical gradients
print('Symbolic:', bce_loss_grad(y, y_hat))

def num_gradient_bce(y, y_hat, epsilon):
return (bce_loss(y, y_hat + 0.5*epsilon) - bce_loss(y, y_hat - 0.5*epsilon))/epsilon

print('Finite differences', num_gradient_bce(y, y_hat, 0.01))

y: [1 1 1 0 1]
y_hat: [0.830152 0.630334 0.153277 0.789648 0.282017]
bceloss(y, y_hat): [0.186147 0.461506 1.875507 1.558971 1.265789]

Symbolic: [-1.204599 -1.586461 -6.524125  4.753925 -3.54589 ]
Finite differences [-1.204614 -1.586495 -6.52644  4.754821 -3.546262]
```


Linear layer

Finally, we need to compute the gradients for the linear layer in our network.

Define a linear model $\mathbf{y} = \mathbf{x}\mathbf{W} + \mathbf{b}$, where

- \mathbf{x} is an input vector of shape N ,
- \mathbf{W} is a weight matrix of shape $N \times M$,
- \mathbf{b} is a bias vector of shape M ,
- \mathbf{y} is the output vector of shape M .

Derive the gradients for \mathbf{y} with respect to the input \mathbf{x} and the parameters \mathbf{W} and \mathbf{b} :

TODO double check this

$$\begin{aligned}\nabla_{\mathbf{x}}\mathbf{y} &= \nabla_{\mathbf{x}}(\mathbf{x}\mathbf{W} + \mathbf{b}) \\ &= \nabla_{\mathbf{x}}(\mathbf{x}\mathbf{W}) + \nabla_{\mathbf{x}}\mathbf{b} \\ &= \mathbf{W}^T + 0 = \mathbf{W}^T \\ \nabla_{\mathbf{W}}\mathbf{y} &= \nabla_{\mathbf{W}}(\mathbf{x}\mathbf{W} + \mathbf{b}) \\ &= \nabla_{\mathbf{W}}(\mathbf{x}\mathbf{W}) + \nabla_{\mathbf{W}}\mathbf{b} \\ &= \mathbf{x}^T + 0 = \mathbf{x}^T \\ \nabla_{\mathbf{b}}\mathbf{y} &= \nabla_{\mathbf{b}}(\mathbf{x}\mathbf{W} + \mathbf{b}) \\ &= \nabla_{\mathbf{b}}(\mathbf{x}\mathbf{W}) + \nabla_{\mathbf{b}}\mathbf{b} \\ &= 0 + I = I\end{aligned}$$

Given the gradient $\nabla_{\mathbf{y}}\mathcal{L}$ for the loss w.r.t. \mathbf{y} , use the chain rule to derive the gradients for the loss w.r.t. \mathbf{x} , \mathbf{W} and \mathbf{b} :

$$\begin{aligned}\nabla_{\mathbf{x}}\mathcal{L} &= \frac{\partial \mathcal{L}}{\partial \mathbf{x}} \\ &= \frac{\partial \mathcal{L}}{\partial \mathbf{y}} \frac{\partial \mathbf{y}}{\partial \mathbf{x}} = \nabla_{\mathbf{y}}\mathcal{L} \nabla_{\mathbf{x}}\mathbf{y} = \nabla_{\mathbf{y}}\mathcal{L} \mathbf{W}^T \\ \nabla_{\mathbf{W}}\mathcal{L} &= \frac{\partial \mathcal{L}}{\partial \mathbf{W}} \\ &= \frac{\partial \mathcal{L}}{\partial \mathbf{y}} \frac{\partial \mathbf{y}}{\partial \mathbf{W}} = \frac{\partial \mathbf{y}}{\partial \mathbf{W}} \frac{\partial \mathcal{L}}{\partial \mathbf{y}} = \nabla_{\mathbf{W}}\mathbf{y} \nabla_{\mathbf{y}}\mathcal{L} = \mathbf{x}^T \nabla_{\mathbf{y}}\mathcal{L} \\ \nabla_{\mathbf{b}}\mathcal{L} &= \frac{\partial \mathcal{L}}{\partial \mathbf{b}} \\ &= \frac{\partial \mathcal{L}}{\partial \mathbf{y}} \frac{\partial \mathbf{y}}{\partial \mathbf{b}} = \nabla_{\mathbf{y}}\mathcal{L} \nabla_{\mathbf{b}}\mathbf{y} = \nabla_{\mathbf{y}}\mathcal{L} I = \nabla_{\mathbf{y}}\mathcal{L}\end{aligned}$$

1.3 Implement a one-layer model

We can now implement a simple one-layer model with a sigmoid activation:

1. Given an input vector \mathbf{x} , weight vector \mathbf{w} and bias b , compute the output \hat{y} :

$$\begin{aligned}h &= \mathbf{x}^T \mathbf{w} + b \\ \hat{y} &= \sigma(h)\end{aligned}$$

1. Compute the BCE loss comparing the prediction \hat{y} with the ground-truth label y .
2. Compute the gradient for the BCE loss and back-propagate this to get the gradient of \mathcal{L} w.r.t. \mathbf{x} :

$$\begin{aligned}\nabla_{\mathbf{x}}\mathcal{L} &= \frac{\partial \mathcal{L}}{\partial \mathbf{x}} \\ &= \frac{\partial \mathcal{L}}{\partial \mathbf{y}} \frac{\partial \mathbf{y}}{\partial \mathbf{x}} \\ &= \frac{\partial \mathcal{L}}{\partial \mathbf{y}} \frac{\partial \mathbf{y}}{\partial \mathbf{h}} \frac{\partial \mathbf{h}}{\partial \mathbf{x}}\end{aligned}$$

Complete the implementation below:

```
In [15]: # initialize parameters
w = np.random.uniform(size=5)
b = np.random.rand()

# implement the model
def fn(x, y):
    # DONE: forward: compute h, y_hat, loss
    h = x.dot(w) + b
    y_hat = sigmoid(h)
    loss = bce_loss(y, y_hat)

    # DONE: backward: compute grad_y_hat, grad_h, grad_x
    grad_y_hat = bce_loss_grad(y, y_hat) # Gradient of loss func w.r.t y_hat
    grad_h = grad_y_hat * sigmoid_grad(h) # Gradient of loss func w.r.t h
    grad_x = grad_h * w # Gradient of loss func w.r.t. x

    return loss, grad_x

# test with a random input
x = np.ones((5,)) # np.random.uniform(size=5)
y = 1

loss, grad_x = fn(x, y)
print("Loss", loss)
print("Gradient", grad_x)
```

```
Loss 0.061126062909673845
Gradient [-0.051182 -0.017641 -0.051237 -0.01718  -0.000469]
```

Use the finite-difference method to check the gradient $\nabla_x \mathcal{L}$:

```
In [16]: # start with some random inputs
x = np.random.uniform(size=5)
y = 1

# set epsilon to a small value
eps = 0.00001

numerical_grad = np.zeros(x.shape)
# compute the gradient for each element of x separately
for i in range(len(x)):
    # compute inputs at -eps/2 and +eps/2
    x_a, x_b = x.copy(), x.copy()
    x_a[i] += eps / 2
    x_b[i] -= eps / 2

    # compute the gradient for this element
    loss_a, _ = fn(x_a, y)
    loss_b, _ = fn(x_b, y)
    numerical_grad[i] = (loss_a - loss_b) / eps

# compute the symbolic gradient
loss, symbolic_grad = fn(x, y)

print("Symbolic gradient")
print(symbolic_grad)
print("Numerical gradient")
print(numerical_grad)
```

```
Symbolic gradient
[-0.205776 -0.070926 -0.205999 -0.069071 -0.001886]
Numerical gradient
[-0.205776 -0.070926 -0.205999 -0.069071 -0.001886]
```

1.4 Implement a linear layer and the sigmoid and ReLU activation functions

We will now construct a simple neural network. We need to implement the following objects:

- **Linear**: a layer that computes $y = x \cdot W + b$.
- **Sigmoid**: a layer that computes $y = \text{sigmoid}(x)$.
- **ReLU**: a layer that computes $y = \text{relu}(x)$.

For each layer class, we need to implement the following methods:

- **forward**: The forward pass that computes the output y given x .
- **backward**: The backward pass that receives the gradient for y and computes the gradients for the input x and the parameters of the layer.
- **step**: The update step that applies the gradient updates to the weights, based on the gradient computed by **backward**.

Implement a class **Linear** that computes $y = x \cdot W + b$:

```

In [17]: # Computes  $y = x * w + b$ .
class Linear:
    def __init__(self, n_in, n_out):
        # initialize the weights randomly,
        # using the Xavier initialization rule for scale
        a = np.sqrt(6 / (n_in * n_out))
        self.W = np.random.uniform(-a, a, size=(n_in, n_out))
        self.b = np.zeros((n_out,))

    def forward(self, x):
        # DONE: compute the forward pass
        y = x.dot(self.W) + self.b
        #print("input of shape", x.shape)
        #print("weights of shape", self.W.shape)
        #print("output of shape", y.shape)
        return y

    def backward(self, x, dy):
        # DONE: compute the backward pass,
        # given dy, compute the gradients for x, W and b
        dx = dy.dot(self.W.T)
        self.dW = x.T.dot(dy)
        self.db = np.sum(dy, axis=0) # np.ones((1, x.shape[0])).dot(dy)
        #print("dy", dy)
        #print(self.db)
        #print(dy.shape)
        #print(np.ones((self.b.shape[0], self.b.shape[0])).shape)
        #print(self.db.shape)
        #print("dx", dx.shape)
        return dx

    def step(self, step):
        #raise NotImplementedError # DONE?
        # TODO: apply a gradient descent update step
        self.W = self.W - step*self.dW # DONE?
        self.b = self.b - step*self.db # DONE?

    def __str__(self):
        return 'Linear %dx%d' % self.W.shape

# Try the new class with some random values.
# Debugging tip: always choose a unique length for each
# dimension, so you'll get an error if you mix them up.
x = np.random.uniform(size=(3, 5))

layer = Linear(5, 7)
y = layer.forward(x)
dx = layer.backward(x, np.ones_like(y))
print('y:', y)
print('dx:', dx)

y: [[-0.089755 -0.206227 -0.413744  0.405997 -0.024196  0.341687  0.419334]
     [-0.137369 -0.164599 -0.416627  0.246486  0.222668  0.199937  0.497592]
     [-0.35406  -0.237123 -0.134597  0.461111  0.103766  0.359734  0.123361]]
dx: [[ 0.066349  0.307645  0.32633  0.11034 -0.092874]
     [ 0.066349  0.307645  0.32633  0.11034 -0.092874]
     [ 0.066349  0.307645  0.32633  0.11034 -0.092874]]

```

Implement a class **Sigmoid** that computes $y = 1 / (1 + \exp(-x))$:

```
In [18]: # Computes  $y = 1 / (1 + \exp(-x))$ .
class Sigmoid:
    def forward(self, x):
        # DONE: compute the forward pass
        return sigmoid(x)

    def backward(self, x, dy):
        # DONE: compute the backward pass,
        # return the gradient for x given dy
        return dy * sigmoid_grad(x)

    def step(self, step_size):
        #raise NotImplementedError # TODO
        return

    def __str__(self):
        return 'Sigmoid'

# try the new class with some random values
x = np.random.uniform(size=(3, 5))

layer = Sigmoid()
y = layer.forward(x)
dx = layer.backward(x, np.ones_like(y))
print('y:', y)
print('dx:', dx)

y: [[0.670335 0.722953 0.587699 0.528931 0.714446]
     [0.70388  0.561949 0.69742  0.673355 0.503916]
     [0.660563 0.52696  0.574788 0.658705 0.532889]]
dx: [[0.220986 0.200292 0.242309 0.249163 0.204013]
     [0.208433 0.246162 0.211025 0.219948 0.249985]
     [0.22422  0.249273 0.244407 0.224813 0.248918]]
```

Implement a class **ReLU** that computes $y = \max(0, x)$:

```
In [19]: # Computes  $y = \max(0, x)$ .
class ReLU:
    def forward(self, x):
        # DONE: compute the forward pass
        return relu(x)

    def backward(self, x, dy):
        # DONE: compute the backward pass,
        # return the gradient for x given dy
        return dy * relu_grad(x)

    def step(self, step_size):
        #raise NotImplementedError # TODO
        return

    def __str__(self):
        return 'ReLU'

# try the new class with some random values
x = np.random.uniform(-10, 10, size=(3, 5))

layer = ReLU()
y = layer.forward(x)
dx = layer.backward(x, np.ones_like(y))
print('y:', y)
print('dx:', dx)

y: [[0.         0.         0.         6.760611 0.        ]
     [4.850774 0.         5.604991 0.         0.        ]
     [8.274109 6.895369 0.         9.71315  3.297718]]
dx: [[0. 0. 0. 1. 0.]
     [1. 0. 1. 0. 0.]
     [1. 1. 0. 1. 1.]]
```

Verify the gradients (using provided code)

The code below will check your implementations using SciPy's finite difference implementation `check_grad` (https://docs.scipy.org/doc/scipy/reference/generated/scipy.optimize.check_grad.html). This is similar to what we did manually before, but automates some of the work.

Run the code and check that the error is not too large.

```
In [20]: ## Verify gradient computations for Linear
# test for dx
layer = Linear(5, 7)
def test_fn(x):
    x = x.reshape(3, 5)
    # multiply the output with a constant to check if
    # the gradient uses dy
    return 2 * np.sum(layer.forward(x))
def test_fn_grad(x):
    x = x.reshape(3, 5)
    # multiply the incoming dy gradient with a constant
    return layer.backward(x, 2 * np.ones((3, 7))).flatten()

err = scipy.optimize.check_grad(test_fn, test_fn_grad,
                                np.random.uniform(-10, 10, size=3 * 5))
print("err on dx:", "OK" if np.abs(err) < 1e-5 else "ERROR", err)

# test for dW
x = np.random.uniform(size=(3, 5))
layer = Linear(5, 7)
def test_fn(w):
    layer.W = w.reshape(5, 7)
    # multiply the output with a constant to check if
    # the gradient uses dy
    return 2 * np.sum(layer.forward(x))
def test_fn_grad(w):
    layer.W = w.reshape(5, 7)
    # multiply the incoming dy gradient with a constant
    layer.backward(x, 2 * np.ones((3, 7)))
    return layer.dW.flatten()

err = scipy.optimize.check_grad(test_fn, test_fn_grad,
                                np.random.uniform(-10, 10, size=5 * 7))
print("err on dW:", "OK" if np.abs(err) < 1e-5 else "ERROR", err)

# test for db
x = np.random.uniform(size=(3, 5,))
layer = Linear(5, 7)
def test_fn(b):
    layer.b = b
    # multiply the output with a constant to check if
    # the gradient uses dy
    return 2 * np.sum(layer.forward(x))
def test_fn_grad(b):
    layer.b = b
    # multiply the incoming dy gradient with a constant
    layer.backward(x, 2 * np.ones((x.shape[0], 7)))
    return layer.db

err = scipy.optimize.check_grad(test_fn, test_fn_grad,
                                np.random.uniform(-10, 10, size=7))
print("err on db:", "OK" if np.abs(err) < 1e-5 else "ERROR", err)

err on dx: OK 8.163466191454083e-07
err on dW: OK 3.0827970880318133e-06
err on db: OK 0.0
```

```
In [21]: ## Verify gradient computation for Sigmoid
# test for dx
layer = Sigmoid()
def test_fn(x):
    # multiply the output with a constant to check if
    # the gradient uses dy
    return np.sum(2 * layer.forward(x))
def test_fn_grad(x):
    # multiply the incoming dy gradient with a constant
    return layer.backward(x, 2 * np.ones(x.shape))

err = scipy.optimize.check_grad(test_fn, test_fn_grad,
                                np.random.uniform(-10, 10, size=5))
print("err on dx:", "OK" if np.abs(err) < 1e-5 else "ERROR", err)

err on dx: OK 2.851954923176839e-08
```

```
In [22]: ## Verify gradient computation for ReLU
# test for dx
layer = ReLU()
def test_fn(x):
    # multiply the output with a constant to check if
    # the gradient uses dy
    return 2 * np.sum(layer.forward(x))
def test_fn_grad(x):
    # multiply the incoming dy gradient with a constant
    return layer.backward(x, 2 * np.ones(x.shape))

err = scipy.optimize.check_grad(test_fn, test_fn_grad,
                                np.random.uniform(1, 10, size=5))
print("err on dx:", "OK" if np.abs(err) < 1e-5 else "ERROR", err)

err on dx: OK 0.0
```

1.5 Construct a neural network with back-propagation

We will use the following container class to implement the network:

1. The `forward` pass computes the output of each layer. We store the intermediate inputs for the backward pass.
2. The `backward` pass computes the gradients for each layer, in reverse order, by using the original input `x` and the gradient `dy` from the previous layer.
3. The `step` function will ask each layer to apply the gradient descent updates to its weights.

Read the code below:

```
In [23]: class Net:
def __init__(self, layers):
    self.layers = layers

def forward(self, x):
    # compute the forward pass for each layer
    trace = []
    for layer in self.layers:
        # compute the forward pass
        y = layer.forward(x)
        # store the original input for the backward pass
        trace.append((layer, x))
        x = y
    # return the final output and the history trace
    return y, trace

def backward(self, trace, dy):
    # compute the backward pass for each layer
    for layer, x in trace[::-1]:
        # compute the backward pass using the original input x
        dy = layer.backward(x, dy)

def step(self, learning_rate):
    # apply the gradient descent updates of each layer
    for layer in self.layers:
        layer.step(learning_rate)

def __str__(self):
    return '\n'.join(str(l) for l in self.layers)
```

1.6 Training the network

We load a simple dataset with 360 handwritten digits.

Each sample has 8×8 pixels, arranged as a 1D vector of 64 features.

We create a binary classification problem with the label 0 for the digits 0 to 4, and 1 for the digits 5 to 9.

```
In [24]: # load the first two classes of the digits dataset
dataset = sklearn.datasets.load_digits()
digits_x, digits_y = dataset['data'], dataset['target']

# create a binary classification problem
digits_y = (digits_y < 5).astype(float)

# plot some of the digits
plt.figure(figsize=(10, 2))
plt.imshow(np.hstack([digits_x[i].reshape(8, 8) for i in range(10)]), cmap='gray')
plt.grid(False)
plt.tight_layout()
plt.axis('off')

# normalize the values to [0, 1]
digits_x -= np.mean(digits_x)
digits_x /= np.std(digits_x)

# print some statistics
print('digits_x.shape:', digits_x.shape)
print('digits_y.shape:', digits_y.shape)
print('min, max values:', np.min(digits_x), np.max(digits_x))
print('labels:', np.unique(digits_y))

digits_x.shape: (1797, 64)
digits_y.shape: (1797,)
min, max values: -0.8117561971974786 1.847470154168513
labels: [0. 1.]
```



We divide the dataset in a train and a test set.


```
In [25]: # make a 50%/50% train/test split
train_prop = 0.5
n_train = int(digits_x.shape[0] * train_prop)

# shuffle the images
idxs = np.random.permutation(digits_x.shape[0])

# take a subset
x = {'train': digits_x[idxs[:n_train]],
     'test':  digits_x[idxs[n_train:]]}
y = {'train': digits_y[idxs[:n_train]],
     'test':  digits_y[idxs[n_train:]]}

print('Training samples:', x['train'].shape[0])
print('Test samples:', x['test'].shape[0])
```

```
Training samples: 898
Test samples: 899
```

We will now implement a function that trains the network. For each epoch, it loops over all minibatches in the training set and updates the network weights. It will then compute the loss and accuracy for the test samples. Finally, it will plot the learning curves.

Read through the code below.

```
In [26]: def fit(net, x, y, epochs=25, learning_rate=0.001, mb_size=10):
# initialize the loss and accuracy history
loss_hist = {'train': [], 'test': []}
accuracy_hist = {'train': [], 'test': []}

for epoch in range(epochs):
# initialize the loss and accuracy for this epoch
loss = {'train': 0.0, 'test': 0.0}
accuracy = {'train': 0.0, 'test': 0.0}

# first train on training data, then evaluate on the test data
for phase in ('train', 'test'):
# compute the number of minibatches
steps = x[phase].shape[0] // mb_size

# loop over all minibatches
for step in range(steps):
# get the samples for the current minibatch
x_mb = x[phase][(step * mb_size):((step + 1) * mb_size)]
y_mb = y[phase][(step * mb_size):((step + 1) * mb_size), None]

# compute the forward pass through the network
pred_y, trace = net.forward(x_mb)

# compute the current loss and accuracy
loss[phase] += np.mean(bce_loss(y_mb, pred_y))
accuracy[phase] += np.mean((y_mb > 0.5) == (pred_y > 0.5))

# only update the network in the training phase
if phase == 'train':
# compute the gradient for the loss
dy = bce_loss_grad(y_mb, pred_y)

# backpropagate the gradient through the network
net.backward(trace, dy)

# update the weights
net.step(learning_rate)

# compute the mean loss and accuracy over all minibatches
loss[phase] = loss[phase] / steps
accuracy[phase] = accuracy[phase] / steps

# add statistics to history
loss_hist[phase].append(loss[phase])
accuracy_hist[phase].append(accuracy[phase])

print('Epoch %3d: loss[train]=%7.4f accuracy[train]=%7.4f loss[test]=%7.4f accuracy[test]=%7.4f' %
      (epoch, loss['train'], accuracy['train'], loss['test'], accuracy['test']))

# plot the learning curves
plt.figure(figsize=(20, 5))

plt.subplot(1, 2, 1)
for phase in loss_hist:
plt.plot(loss_hist[phase], label=phase)
plt.title('BCE loss')
plt.xlabel('Epoch')
plt.legend()

plt.subplot(1, 2, 2)
for phase in accuracy_hist:
plt.plot(accuracy_hist[phase], label=phase)
plt.title('Accuracy')
plt.xlabel('Epoch')
plt.legend()
```

We will define a two-layer network:


- A linear layer that maps the 64 features of the input to 32 features.
- A ReLU activation function.
- A linear layer that maps the 32 features to the 1 output features.
- A sigmoid activation function that maps the output to [0, 1].

Train the network and inspect the results. Tune the hyperparameters to get a good result.

```
In [27]: # construct network
net = Net([
    Linear(64, 32),
    ReLU(),
    Linear(32, 1),
    Sigmoid())])

# TODO: tune the hyperparameters
fit(net, x, y,
    epochs = 25,
    learning_rate = 0.01,
    mb_size = 10)
```

Epoch	0:	loss[train]= 0.4225	accuracy[train]= 0.8213	loss[test]= 0.3472	accuracy[test]= 0.8584
Epoch	1:	loss[train]= 0.2390	accuracy[train]= 0.9124	loss[test]= 0.3097	accuracy[test]= 0.8719
Epoch	2:	loss[train]= 0.1773	accuracy[train]= 0.9393	loss[test]= 0.2393	accuracy[test]= 0.9101
Epoch	3:	loss[train]= 0.1374	accuracy[train]= 0.9607	loss[test]= 0.1878	accuracy[test]= 0.9292
Epoch	4:	loss[train]= 0.1103	accuracy[train]= 0.9663	loss[test]= 0.1622	accuracy[test]= 0.9393
Epoch	5:	loss[train]= 0.0908	accuracy[train]= 0.9730	loss[test]= 0.1405	accuracy[test]= 0.9449
Epoch	6:	loss[train]= 0.0765	accuracy[train]= 0.9787	loss[test]= 0.1318	accuracy[test]= 0.9472
Epoch	7:	loss[train]= 0.0649	accuracy[train]= 0.9798	loss[test]= 0.1207	accuracy[test]= 0.9551
Epoch	8:	loss[train]= 0.0558	accuracy[train]= 0.9843	loss[test]= 0.1163	accuracy[test]= 0.9573
Epoch	9:	loss[train]= 0.0476	accuracy[train]= 0.9899	loss[test]= 0.1143	accuracy[test]= 0.9573
Epoch	10:	loss[train]= 0.0412	accuracy[train]= 0.9910	loss[test]= 0.1124	accuracy[test]= 0.9607
Epoch	11:	loss[train]= 0.0357	accuracy[train]= 0.9921	loss[test]= 0.1127	accuracy[test]= 0.9596
Epoch	12:	loss[train]= 0.0314	accuracy[train]= 0.9933	loss[test]= 0.1101	accuracy[test]= 0.9618
Epoch	13:	loss[train]= 0.0277	accuracy[train]= 0.9944	loss[test]= 0.1095	accuracy[test]= 0.9618
Epoch	14:	loss[train]= 0.0246	accuracy[train]= 0.9955	loss[test]= 0.1073	accuracy[test]= 0.9618
Epoch	15:	loss[train]= 0.0215	accuracy[train]= 0.9966	loss[test]= 0.1074	accuracy[test]= 0.9607
Epoch	16:	loss[train]= 0.0192	accuracy[train]= 0.9966	loss[test]= 0.1054	accuracy[test]= 0.9618
Epoch	17:	loss[train]= 0.0173	accuracy[train]= 0.9978	loss[test]= 0.1035	accuracy[test]= 0.9618
Epoch	18:	loss[train]= 0.0152	accuracy[train]= 0.9989	loss[test]= 0.1055	accuracy[test]= 0.9629
Epoch	19:	loss[train]= 0.0139	accuracy[train]= 0.9989	loss[test]= 0.1042	accuracy[test]= 0.9652
Epoch	20:	loss[train]= 0.0127	accuracy[train]= 1.0000	loss[test]= 0.1030	accuracy[test]= 0.9652
Epoch	21:	loss[train]= 0.0113	accuracy[train]= 1.0000	loss[test]= 0.1035	accuracy[test]= 0.9663
Epoch	22:	loss[train]= 0.0104	accuracy[train]= 1.0000	loss[test]= 0.1029	accuracy[test]= 0.9663
Epoch	23:	loss[train]= 0.0095	accuracy[train]= 1.0000	loss[test]= 0.1046	accuracy[test]= 0.9652
Epoch	24:	loss[train]= 0.0087	accuracy[train]= 1.0000	loss[test]= 0.1039	accuracy[test]= 0.9663



Which of the hyperparameters (number of epochs, learning rate, minibatch size) was most important? How did they influence your results?

Changing the minibatch size seems to make little difference. Though, if the learning rate is failry high, larger minibatches seem to lead to more fluctuation in accuracy over time. The learning rate makes quite the difference and if we define it too high, it runs into problems as it tries to divide by zero. The number of epochs only makes a difference in as far as it gets enough epochs to reach 100% accuracy on the training set.


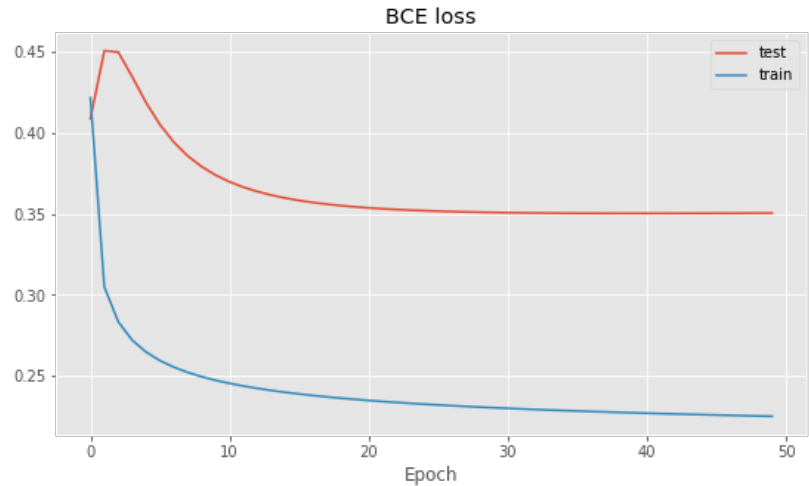
Repeat the experiment with a the same network, but remove the ReLU activation in the middle: [Linear, Linear, Sigmoid] .

In [28]:

```
# TODO: Your code here.
# construct network
net2 = Net([
    Linear(64, 32),
    Linear(32, 1),
    Sigmoid())

# TODO: tune the hyperparameters
fit(net2, x, y,
    epochs = 50,
    learning_rate = 0.01,
    mb_size = 10)
```

Epoch	0:	loss[train]= 0.4218	accuracy[train]= 0.8056	loss[test]= 0.4088	accuracy[test]= 0.8213
Epoch	1:	loss[train]= 0.3050	accuracy[train]= 0.8831	loss[test]= 0.4508	accuracy[test]= 0.8135
Epoch	2:	loss[train]= 0.2835	accuracy[train]= 0.8910	loss[test]= 0.4499	accuracy[test]= 0.8124
Epoch	3:	loss[train]= 0.2722	accuracy[train]= 0.8944	loss[test]= 0.4348	accuracy[test]= 0.8225
Epoch	4:	loss[train]= 0.2648	accuracy[train]= 0.8944	loss[test]= 0.4188	accuracy[test]= 0.8292
Epoch	5:	loss[train]= 0.2595	accuracy[train]= 0.8978	loss[test]= 0.4052	accuracy[test]= 0.8337
Epoch	6:	loss[train]= 0.2555	accuracy[train]= 0.9000	loss[test]= 0.3943	accuracy[test]= 0.8416
Epoch	7:	loss[train]= 0.2522	accuracy[train]= 0.9011	loss[test]= 0.3858	accuracy[test]= 0.8472
Epoch	8:	loss[train]= 0.2496	accuracy[train]= 0.9034	loss[test]= 0.3792	accuracy[test]= 0.8528
Epoch	9:	loss[train]= 0.2474	accuracy[train]= 0.9045	loss[test]= 0.3740	accuracy[test]= 0.8573
Epoch	10:	loss[train]= 0.2455	accuracy[train]= 0.9056	loss[test]= 0.3699	accuracy[test]= 0.8551
Epoch	11:	loss[train]= 0.2438	accuracy[train]= 0.9067	loss[test]= 0.3665	accuracy[test]= 0.8551
Epoch	12:	loss[train]= 0.2424	accuracy[train]= 0.9056	loss[test]= 0.3639	accuracy[test]= 0.8584
Epoch	13:	loss[train]= 0.2411	accuracy[train]= 0.9056	loss[test]= 0.3617	accuracy[test]= 0.8607
Epoch	14:	loss[train]= 0.2400	accuracy[train]= 0.9056	loss[test]= 0.3599	accuracy[test]= 0.8607
Epoch	15:	loss[train]= 0.2389	accuracy[train]= 0.9067	loss[test]= 0.3584	accuracy[test]= 0.8607
Epoch	16:	loss[train]= 0.2380	accuracy[train]= 0.9067	loss[test]= 0.3571	accuracy[test]= 0.8629
Epoch	17:	loss[train]= 0.2371	accuracy[train]= 0.9067	loss[test]= 0.3561	accuracy[test]= 0.8618
Epoch	18:	loss[train]= 0.2363	accuracy[train]= 0.9079	loss[test]= 0.3552	accuracy[test]= 0.8607
Epoch	19:	loss[train]= 0.2356	accuracy[train]= 0.9079	loss[test]= 0.3544	accuracy[test]= 0.8618
Epoch	20:	loss[train]= 0.2349	accuracy[train]= 0.9079	loss[test]= 0.3538	accuracy[test]= 0.8629
Epoch	21:	loss[train]= 0.2343	accuracy[train]= 0.9079	loss[test]= 0.3532	accuracy[test]= 0.8629
Epoch	22:	loss[train]= 0.2337	accuracy[train]= 0.9067	loss[test]= 0.3528	accuracy[test]= 0.8629
Epoch	23:	loss[train]= 0.2332	accuracy[train]= 0.9067	loss[test]= 0.3524	accuracy[test]= 0.8629
Epoch	24:	loss[train]= 0.2326	accuracy[train]= 0.9056	loss[test]= 0.3520	accuracy[test]= 0.8629
Epoch	25:	loss[train]= 0.2322	accuracy[train]= 0.9067	loss[test]= 0.3517	accuracy[test]= 0.8652
Epoch	26:	loss[train]= 0.2317	accuracy[train]= 0.9079	loss[test]= 0.3515	accuracy[test]= 0.8652
Epoch	27:	loss[train]= 0.2313	accuracy[train]= 0.9079	loss[test]= 0.3513	accuracy[test]= 0.8652
Epoch	28:	loss[train]= 0.2308	accuracy[train]= 0.9079	loss[test]= 0.3511	accuracy[test]= 0.8652
Epoch	29:	loss[train]= 0.2305	accuracy[train]= 0.9079	loss[test]= 0.3510	accuracy[test]= 0.8652
Epoch	30:	loss[train]= 0.2301	accuracy[train]= 0.9079	loss[test]= 0.3508	accuracy[test]= 0.8652
Epoch	31:	loss[train]= 0.2297	accuracy[train]= 0.9079	loss[test]= 0.3507	accuracy[test]= 0.8663
Epoch	32:	loss[train]= 0.2294	accuracy[train]= 0.9079	loss[test]= 0.3506	accuracy[test]= 0.8652
Epoch	33:	loss[train]= 0.2290	accuracy[train]= 0.9079	loss[test]= 0.3506	accuracy[test]= 0.8629
Epoch	34:	loss[train]= 0.2287	accuracy[train]= 0.9079	loss[test]= 0.3505	accuracy[test]= 0.8629
Epoch	35:	loss[train]= 0.2284	accuracy[train]= 0.9079	loss[test]= 0.3505	accuracy[test]= 0.8629
Epoch	36:	loss[train]= 0.2281	accuracy[train]= 0.9079	loss[test]= 0.3504	accuracy[test]= 0.8629
Epoch	37:	loss[train]= 0.2279	accuracy[train]= 0.9079	loss[test]= 0.3504	accuracy[test]= 0.8629
Epoch	38:	loss[train]= 0.2276	accuracy[train]= 0.9090	loss[test]= 0.3504	accuracy[test]= 0.8629
Epoch	39:	loss[train]= 0.2273	accuracy[train]= 0.9090	loss[test]= 0.3504	accuracy[test]= 0.8629
Epoch	40:	loss[train]= 0.2271	accuracy[train]= 0.9090	loss[test]= 0.3504	accuracy[test]= 0.8652
Epoch	41:	loss[train]= 0.2268	accuracy[train]= 0.9090	loss[test]= 0.3504	accuracy[test]= 0.8652
Epoch	42:	loss[train]= 0.2266	accuracy[train]= 0.9090	loss[test]= 0.3504	accuracy[test]= 0.8652
Epoch	43:	loss[train]= 0.2264	accuracy[train]= 0.9090	loss[test]= 0.3504	accuracy[test]= 0.8652
Epoch	44:	loss[train]= 0.2262	accuracy[train]= 0.9090	loss[test]= 0.3505	accuracy[test]= 0.8663
Epoch	45:	loss[train]= 0.2260	accuracy[train]= 0.9090	loss[test]= 0.3505	accuracy[test]= 0.8663
Epoch	46:	loss[train]= 0.2257	accuracy[train]= 0.9090	loss[test]= 0.3505	accuracy[test]= 0.8663
Epoch	47:	loss[train]= 0.2256	accuracy[train]= 0.9101	loss[test]= 0.3505	accuracy[test]= 0.8652
Epoch	48:	loss[train]= 0.2254	accuracy[train]= 0.9090	loss[test]= 0.3506	accuracy[test]= 0.8663
Epoch	49:	loss[train]= 0.2252	accuracy[train]= 0.9090	loss[test]= 0.3506	accuracy[test]= 0.8674



How does the performance of this network compare with the previous network. Can you explain this result? How does removing the ReLU affect the model?

After the model minimizes the loss on the test set, it doesn't slowly improve anymore like the other model did, but just stays the same, while the previous network continued to improve performance on the test set after the "bend" in the graph.

Similarly, the accuracy on the test set stays stable over the entire training, suggesting almost no actual improvement is taking place.

Even for a large amount of epochs, the training accuracy is quite slow to, and never quite reaches 1.00, meaning that it is almost not even powerful enough to overfit well.

Removing the ReLU means that two linear layers follow one another, which is equivalent to a single linear layer in terms of the decision boundaries it can represent.

In a sense this network just has a lot less descriptive power / is less able to model distributions that aren't very simple.

Create a network with one linear layer followed by a sigmoid activation:



```
net = Net([Linear(...), Sigmoid()])
```

Train this network. Compare the results with the [Linear, ReLU, Linear, Sigmoid] and [Linear, Linear, Sigmoid] networks you trained before, and explain the results.

```
In [29]: # TODO: Your code here.
# construct network
net3 = Net([
    Linear(64, 1),
    Sigmoid()])

# TODO: tune the hyperparameters
fit(net3, x, y,
    epochs = 25,
    learning_rate = 0.01,
    mb_size = 10)
```

Epoch	0:	loss[train]= 0.4821	accuracy[train]= 0.7719	loss[test]= 0.4286	accuracy[test]= 0.8135
Epoch	1:	loss[train]= 0.3209	accuracy[train]= 0.8730	loss[test]= 0.3943	accuracy[test]= 0.8348
Epoch	2:	loss[train]= 0.2907	accuracy[train]= 0.8876	loss[test]= 0.3769	accuracy[test]= 0.8494
Epoch	3:	loss[train]= 0.2754	accuracy[train]= 0.8899	loss[test]= 0.3661	accuracy[test]= 0.8573
Epoch	4:	loss[train]= 0.2658	accuracy[train]= 0.8966	loss[test]= 0.3589	accuracy[test]= 0.8607
Epoch	5:	loss[train]= 0.2591	accuracy[train]= 0.8955	loss[test]= 0.3537	accuracy[test]= 0.8573
Epoch	6:	loss[train]= 0.2542	accuracy[train]= 0.8966	loss[test]= 0.3501	accuracy[test]= 0.8640
Epoch	7:	loss[train]= 0.2504	accuracy[train]= 0.9000	loss[test]= 0.3474	accuracy[test]= 0.8640
Epoch	8:	loss[train]= 0.2473	accuracy[train]= 0.9011	loss[test]= 0.3454	accuracy[test]= 0.8663
Epoch	9:	loss[train]= 0.2449	accuracy[train]= 0.9034	loss[test]= 0.3441	accuracy[test]= 0.8685
Epoch	10:	loss[train]= 0.2428	accuracy[train]= 0.9056	loss[test]= 0.3431	accuracy[test]= 0.8674
Epoch	11:	loss[train]= 0.2411	accuracy[train]= 0.9079	loss[test]= 0.3424	accuracy[test]= 0.8685
Epoch	12:	loss[train]= 0.2396	accuracy[train]= 0.9090	loss[test]= 0.3420	accuracy[test]= 0.8697
Epoch	13:	loss[train]= 0.2383	accuracy[train]= 0.9101	loss[test]= 0.3418	accuracy[test]= 0.8719
Epoch	14:	loss[train]= 0.2371	accuracy[train]= 0.9124	loss[test]= 0.3418	accuracy[test]= 0.8719
Epoch	15:	loss[train]= 0.2361	accuracy[train]= 0.9124	loss[test]= 0.3419	accuracy[test]= 0.8742
Epoch	16:	loss[train]= 0.2352	accuracy[train]= 0.9146	loss[test]= 0.3421	accuracy[test]= 0.8730
Epoch	17:	loss[train]= 0.2344	accuracy[train]= 0.9146	loss[test]= 0.3423	accuracy[test]= 0.8742
Epoch	18:	loss[train]= 0.2337	accuracy[train]= 0.9135	loss[test]= 0.3427	accuracy[test]= 0.8719
Epoch	19:	loss[train]= 0.2330	accuracy[train]= 0.9112	loss[test]= 0.3431	accuracy[test]= 0.8708
Epoch	20:	loss[train]= 0.2324	accuracy[train]= 0.9112	loss[test]= 0.3435	accuracy[test]= 0.8708
Epoch	21:	loss[train]= 0.2318	accuracy[train]= 0.9112	loss[test]= 0.3440	accuracy[test]= 0.8708
Epoch	22:	loss[train]= 0.2313	accuracy[train]= 0.9101	loss[test]= 0.3444	accuracy[test]= 0.8697
Epoch	23:	loss[train]= 0.2308	accuracy[train]= 0.9101	loss[test]= 0.3449	accuracy[test]= 0.8708
Epoch	24:	loss[train]= 0.2304	accuracy[train]= 0.9090	loss[test]= 0.3455	accuracy[test]= 0.8685



Discuss your results.



It does not learn as well as the complete network, but on a similar level as the previous one. This makes sense, as two linear layers behind each other can just be combined into one single linear layer (since linear + linear remains linear). So the extra linear layer does not add much.

Try a deeper network (e.g., four linear layers) to see if this can improve the results further.


```
In [30]: # construct network
net = Net([
    Linear(64, 48),
    ReLU(),
    Linear(48, 32),
    ReLU(),
    Linear(32, 21),
    ReLU(),
    Linear(21, 1),
    Sigmoid())])

# TODO: tune the hyperparameters
fit(net, x, y,
    epochs = 50,
    learning_rate = 0.001,
    mb_size = 10)
```

Epoch	0:	loss[train]= 0.6922	accuracy[train]= 0.5281	loss[test]= 0.6948	accuracy[test]= 0.4663
Epoch	1:	loss[train]= 0.6903	accuracy[train]= 0.5371	loss[test]= 0.6962	accuracy[test]= 0.4663
Epoch	2:	loss[train]= 0.6894	accuracy[train]= 0.5371	loss[test]= 0.6971	accuracy[test]= 0.4663
Epoch	3:	loss[train]= 0.6887	accuracy[train]= 0.5371	loss[test]= 0.6975	accuracy[test]= 0.4663
Epoch	4:	loss[train]= 0.6879	accuracy[train]= 0.5371	loss[test]= 0.6975	accuracy[test]= 0.4663
Epoch	5:	loss[train]= 0.6870	accuracy[train]= 0.5371	loss[test]= 0.6971	accuracy[test]= 0.4663
Epoch	6:	loss[train]= 0.6858	accuracy[train]= 0.5371	loss[test]= 0.6962	accuracy[test]= 0.4663
Epoch	7:	loss[train]= 0.6842	accuracy[train]= 0.5371	loss[test]= 0.6948	accuracy[test]= 0.4663
Epoch	8:	loss[train]= 0.6820	accuracy[train]= 0.5371	loss[test]= 0.6926	accuracy[test]= 0.4663
Epoch	9:	loss[train]= 0.6787	accuracy[train]= 0.5371	loss[test]= 0.6890	accuracy[test]= 0.4663
Epoch	10:	loss[train]= 0.6736	accuracy[train]= 0.5371	loss[test]= 0.6836	accuracy[test]= 0.4663
Epoch	11:	loss[train]= 0.6652	accuracy[train]= 0.5371	loss[test]= 0.6746	accuracy[test]= 0.4663
Epoch	12:	loss[train]= 0.6516	accuracy[train]= 0.5528	loss[test]= 0.6605	accuracy[test]= 0.4663
Epoch	13:	loss[train]= 0.6299	accuracy[train]= 0.6404	loss[test]= 0.6388	accuracy[test]= 0.5955
Epoch	14:	loss[train]= 0.5971	accuracy[train]= 0.7393	loss[test]= 0.6087	accuracy[test]= 0.6562
Epoch	15:	loss[train]= 0.5532	accuracy[train]= 0.7955	loss[test]= 0.5686	accuracy[test]= 0.7270
Epoch	16:	loss[train]= 0.5026	accuracy[train]= 0.8494	loss[test]= 0.5224	accuracy[test]= 0.7854
Epoch	17:	loss[train]= 0.4521	accuracy[train]= 0.8753	loss[test]= 0.4786	accuracy[test]= 0.8169
Epoch	18:	loss[train]= 0.4062	accuracy[train]= 0.8820	loss[test]= 0.4381	accuracy[test]= 0.8382
Epoch	19:	loss[train]= 0.3641	accuracy[train]= 0.8966	loss[test]= 0.3996	accuracy[test]= 0.8472
Epoch	20:	loss[train]= 0.3259	accuracy[train]= 0.9056	loss[test]= 0.3643	accuracy[test]= 0.8652
Epoch	21:	loss[train]= 0.2915	accuracy[train]= 0.9146	loss[test]= 0.3405	accuracy[test]= 0.8764
Epoch	22:	loss[train]= 0.2613	accuracy[train]= 0.9169	loss[test]= 0.3150	accuracy[test]= 0.8854
Epoch	23:	loss[train]= 0.2358	accuracy[train]= 0.9191	loss[test]= 0.2913	accuracy[test]= 0.8966
Epoch	24:	loss[train]= 0.2145	accuracy[train]= 0.9270	loss[test]= 0.2702	accuracy[test]= 0.9034
Epoch	25:	loss[train]= 0.1955	accuracy[train]= 0.9326	loss[test]= 0.2505	accuracy[test]= 0.9101
Epoch	26:	loss[train]= 0.1788	accuracy[train]= 0.9382	loss[test]= 0.2321	accuracy[test]= 0.9146
Epoch	27:	loss[train]= 0.1635	accuracy[train]= 0.9449	loss[test]= 0.2160	accuracy[test]= 0.9247
Epoch	28:	loss[train]= 0.1497	accuracy[train]= 0.9494	loss[test]= 0.2012	accuracy[test]= 0.9315
Epoch	29:	loss[train]= 0.1368	accuracy[train]= 0.9551	loss[test]= 0.1881	accuracy[test]= 0.9382
Epoch	30:	loss[train]= 0.1250	accuracy[train]= 0.9596	loss[test]= 0.1766	accuracy[test]= 0.9438
Epoch	31:	loss[train]= 0.1141	accuracy[train]= 0.9640	loss[test]= 0.1669	accuracy[test]= 0.9461
Epoch	32:	loss[train]= 0.1042	accuracy[train]= 0.9652	loss[test]= 0.1585	accuracy[test]= 0.9461
Epoch	33:	loss[train]= 0.0952	accuracy[train]= 0.9652	loss[test]= 0.1517	accuracy[test]= 0.9483
Epoch	34:	loss[train]= 0.0872	accuracy[train]= 0.9697	loss[test]= 0.1459	accuracy[test]= 0.9506
Epoch	35:	loss[train]= 0.0799	accuracy[train]= 0.9742	loss[test]= 0.1411	accuracy[test]= 0.9528
Epoch	36:	loss[train]= 0.0732	accuracy[train]= 0.9775	loss[test]= 0.1371	accuracy[test]= 0.9539
Epoch	37:	loss[train]= 0.0671	accuracy[train]= 0.9798	loss[test]= 0.1340	accuracy[test]= 0.9562
Epoch	38:	loss[train]= 0.0619	accuracy[train]= 0.9831	loss[test]= 0.1315	accuracy[test]= 0.9562
Epoch	39:	loss[train]= 0.0572	accuracy[train]= 0.9831	loss[test]= 0.1299	accuracy[test]= 0.9573
Epoch	40:	loss[train]= 0.0528	accuracy[train]= 0.9854	loss[test]= 0.1284	accuracy[test]= 0.9573
Epoch	41:	loss[train]= 0.0492	accuracy[train]= 0.9865	loss[test]= 0.1268	accuracy[test]= 0.9562
Epoch	42:	loss[train]= 0.0454	accuracy[train]= 0.9865	loss[test]= 0.1264	accuracy[test]= 0.9573
Epoch	43:	loss[train]= 0.0424	accuracy[train]= 0.9876	loss[test]= 0.1251	accuracy[test]= 0.9562
Epoch	44:	loss[train]= 0.0393	accuracy[train]= 0.9910	loss[test]= 0.1246	accuracy[test]= 0.9562
Epoch	45:	loss[train]= 0.0368	accuracy[train]= 0.9921	loss[test]= 0.1244	accuracy[test]= 0.9573
Epoch	46:	loss[train]= 0.0344	accuracy[train]= 0.9933	loss[test]= 0.1238	accuracy[test]= 0.9573
Epoch	47:	loss[train]= 0.0323	accuracy[train]= 0.9933	loss[test]= 0.1236	accuracy[test]= 0.9584
Epoch	48:	loss[train]= 0.0303	accuracy[train]= 0.9955	loss[test]= 0.1235	accuracy[test]= 0.9573
Epoch	49:	loss[train]= 0.0285	accuracy[train]= 0.9955	loss[test]= 0.1229	accuracy[test]= 0.9584



Discuss your findings. Were you able to obtain a perfect classification? Explain the learning curves.

We almost got a perfect classification, but not quite yet on the test set. We see that at the start the network learns quite a lot, because there is still quite a lot to learn. Later on, there is not much to learn left, and the network learns way more slowly. It is normal that the performance on the test set remains slightly below the training set, as we do not use the test set in the learning procedure. Fortunately, they do stay quite close together, meaning that the network has learned to generalise relatively well.

1.7 Final questions

You now have some experience training neural networks. Time for a few final questions.

What is the influence of the learning rate? What happens if the learning rate is too low or too high?

The learning rate, as the name suggests, determines how fast the network learns. If it is too low, the network will only converge very slowly after many iterations, which can become quite costly. If it is too high, the learning becomes erratic and might even diverge, which means the results can also take longer or no result will be reached at all.

What is the role of the minibatch size in SGD? Explain the downsides of a minibatch size that is too small or too high.

Minibatch size is a tradeoff between following an accurate gradient for a somewhat higher computation cost versus an approximate gradient that is faster to compute, multiple times per epoch. Smaller batches allow you to update weights more often, but this update is less accurate.

In the linear layer, we initialized the weights w with random weights, but we initialized the bias b with zeros. What would happen if the weights w were initialised as zeros? Why is this not a problem for the bias?

If we initialize the weights with zero, multiplying with x will again result in an all zero matrix, losing all information of the previous layers. This will make it hard for the gradient to find a direction / diverge in function from other neurons on that layer. Furthermore, the dx that is passed on to the previous layer is based on the weights, so until the weights on this layer slowly leave zero, the gradient propagated back will be tiny.

For b , this is not a problem, because it is only the bias, and if x and W are not zero, then the network will still learn things and update b . The propagated gradient is not based on b .

The end

Well done! Please double check the instructions at the top before you submit your results.